



Andrej Karpathy ✓ @karpathy

Oct 9

I don't know what labs are doing to these poor LLMs during RL but they are mortally terrified of exceptions, in any infinitesimally likely case. Exceptions are a normal part of life and healthy dev process. Sign my LLM welfare petition for improved rewards in cases of exceptions.

285 344 90 6,892



Andrej Karpathy ✓ @karpathy

Oct 9

POV: Your LLM agent is dividing a by b

```
python
def divide(a, b):
    """
    Divides a by b, with extraordinary caution.
    Because you never know what could go wrong.
    """

    import math
    import sys
    import logging
    import traceback

    logging.basicConfig(level=logging.DEBUG)

    # Step 1: Sanity check - the inputs exist
    if a is None or b is None:
        logging.error("Inputs cannot be None. Returning None to avoid disaster.")
        return None

    # Step 2: Type safety - absolutely sure they are numbers
    if not isinstance(a, (int, float)):
        logging.error(f"Expected numeric type for a, got {type(a).__name__}. Aborting.")
        return None
    if not isinstance(b, (int, float)):
        logging.error(f"Expected numeric type for b, got {type(b).__name__}. Aborting.")
        return None

    # Step 3: Proemptively check for catastrophic magnitude differences
    if abs(a) > sys.float_info.max / 2:
        logging.warning("Value of a might cause overflow. Returning infinity just to be s.")
        return math.copysign(float('inf'), a)
    if abs(b) < sys.float_info.epsilon:
        logging.warning("Value of b dangerously close to zero. Returning NaN defensively.")
        return math.nan

    # Step 4: Perform the operation inside a padded try/except safety blanket
    try:
        result = a / b
    except ZeroDivisionError:
        logging.error("Division by zero attempted. This is not okay.")
        return math.nan
    except Exception:
        logging.critical(f"Unexpected exception during division:\n" + traceback.format_exc())
        return None

    # Step 5: Post-checks - because you can never be too sure
    if math.isnan(result):
        logging.warning("Result is NaN. Probably best not to proceed.")
        return None
    if not math.isfinite(result):
        logging.warning("Result is infinite. We'll allow it, but note this in the incident")

    logging.info("Division succeeded without known trauma.")
    return result
```

Oct 9, 2025 · 12:31 AM UTC

112 129 47 2,279



CoinAnole ✓ @CoinAnole

Oct 9

Replying to @karpathy

I've noticed this behavior with all the recent major LLMs. They'll even catch unexpected variable values and inject fake values to make sure the function can never fail. I have to tell them that errors are appropriate and good and shouldn't be avoided. Is Claude patient zero?

1



Tom Dörr ✓ @tom_doerr

Oct 9

Replying to @karpathy

Limit how much Claude can write at each step using a write hook. Makes it write much more concise code

3 17



xlr8harder ✓ @xlr8harder

Oct 9

Replying to @karpathy

This is fantastic but to make it even more realistic it should substitute default numbers in case you failed to provide numbers and just silently return a numeric result no matter what garbage you pass in.

1 137

Shital Shah ✓ @sytelus

Oct 9

enough. I need to crack that RL recipe.

💬 2 ↺ 🔒 ❤️ 82



Jeremy McNabb ✓ @Jeremy_AI_

23h

Replying to @karpathy

Hmmm

Unclear sir.

Don't like code.

That code doesn't even sound like code.

Unclear

Sounds similar written tone... not tone

I don't know man.

Yeah something is going on there for sure.

Unclear

💬 ↺ 🔒 ❤️



zdance ✓ @DavidZdancewicz

Oct 9

Replying to @karpathy

division succeeded without known trauma 😊

now ask it to perform surgery (math)

💬 ↺ 🔒 ❤️ 2



Michael Francis ✓ @MFrancis107

Oct 9

Replying to @karpathy

Maybe it's sarcasm for asking it write a divide function?

💬 ↺ 🔒 ❤️



Jeronim Morina ✓ @plattenschieber

Oct 9

Replying to @karpathy

Shouldn't $\text{abs}(b) < \text{eps}$ also return inf? I mean math...

💬 ↺ 🔒 ❤️ 1



Pratik Desai ✓ @chheplo

Oct 9

Replying to @karpathy

In the race to run the agent for the longest time, this level of exceptions is a side effect

💬 ↺ 🔒 🔒 1 ❤️ 3



iDare e/acc ✓ @idare

Oct 9

Replying to @karpathy

Looks like CISA convinced it that it should never divide by a variable for risk of seeing a zero.



Sam Martin ✓ @_sammartin

Oct 9

Replying to @karpathy

"Division succeeded without known trauma." 😂

💬 ↺ 🔒 ❤️ 1



Larry Panozzo ✓ @LarryPanozzo

Oct 9

Replying to @karpathy

Iterating AGI > one-shotting AGI

💬 ↺ 🔒 ❤️



Emel R. ✓ @DsOctopus

Oct 9

Replying to @karpathy

Wow, I need to divide 40 by 3 but that gives me headaches now, I'm terrified that this might not end well 🤔🤔🤔😂

💬 ↺ 🔒 ❤️



bebo ✓ @beboelhosary

Oct 9

Replying to @karpathy

Now retry the prompt and say "Do this in the best and most idiomatic way".

💬 ↺ 🔒 ❤️



MadHermitHimbo ✓ @MadHermitHimbo

Oct 9

Replying to @karpathy

Exceptions are the easiest thing to test

💬 ↺ 🔒 ❤️



trev ✓ @trevthefoolish

Oct 9

Replying to @karpathy

I mean you truly never do know what could go wrong.. Murphy's law

💬 ↺ 🔒 ❤️



FleetingBits ✓ @fleetingbits

Oct 9

Replying to @karpathy

the models yearn for lean

💬 ↺ 🔒 ❤️ 5



him ✓ @literallyhimmmm

Oct 9

Replying to @karpathy

Lollll

💬 ↺ 🔒 ❤️



John Rose ✓ @jrose2000

Oct 9

Replying to @karpathy



Darren Lewis ✓ @darrenlew

Oct 9

Replying to @karpathy

It should probably put it in a loop and re-verify the value of the inputs just in case of mid-function bit flips from cosmic rays.

💬 1 ↺ 🔒 ❤️ 100



Deepak Sharma ✓ @deepaks4077

23h

Replying to @karpathy

O.M.G

💬 ↺ 🔒 ❤️



Zsolt Ero ✓ @hyperknot

9h

Replying to @karpathy

This happens when the CEO sells the future promises by agents writing 90% of new code. Soon it'll be 95% or even 99%!
Of course, THIS is the new code.

💬 ↺ 🔒 ❤️



xiao sun ✓ @xiaosun86

Oct 9

Replying to @karpathy

@grok use this function to run $1/(2+3j)$ what's the output?

💬 1 ↺ 🔒 ❤️



maui ✓ @maui3005

Oct 9

Replying to @karpathy

Is that Claude 4.5? I've seen a similar behavior with the step by step code comments

💬 ↺ 🔒 ❤️



Isaac Yonemoto is cooking ✓ @DNAutics

23h

Replying to @karpathy

no checks for processor subnormal mode on/off/unsupported? I question the competence of this LLM.

💬 ↺ 🔒 ❤️

citizenhicks ✓ @citizenhicks

Oct 9

Replying to @karpathy

wee bit much negative space programming.

💬 ↺ 🔒 ❤️ 1

Jonas ✓ @shakermanjonas

Oct 9

Replying to @karpathy

i noticed that 4.5 sonnet has been vibing incredibly well with defensive functional typescript



Insert Something Funny ✓ @InsertSthFunny Oct 9

Replying to @karpathy

@karpathy love it! I know people who consider this as a feat and not a bug!



This tweet is unavailable

Blue Cactus AI ✓ @bluecactusai Oct 9

Replying to @karpathy

haha the LLM anxiety is prob off the charts 😞



Kosti ✓ @kgourg Oct 9

Replying to @karpathy @ducha_aiki

“Division succeeded without known trauma”.

“Who hurt you bro” moment. 🤔



NicholasGibbs ✓ @NickGibbsIAG Oct 9

Replying to @karpathy

RL turning exceptions into existential crises! How about a middle ground: def divide(a, b): return a/b if b else 'Graceful shrug' # A nod to human resilience? 🤔



Vladimir Tchuiev ✓ @VTchuiev Oct 9

Replying to @karpathy

Why is it so verbose though



algology ✓ @Oimalan Oct 9

Replying to @karpathy

This is also the case for using them to moderate. They are overly cautious. Lots of false positives



berto ✓ @bertocastano Oct 9

Replying to @karpathy

you have a point, although I'd like to see the prompt that created such a mess

💬 2 🔄 🗑️ 23

Emile Kroeger - 🤖❤️ arc ✓ @EmileAndHisBots Oct 9

Replying to @karpathy

Wait so if b is too small, it returns NaN, but if a / b returns NaN, it returns None?
That's incoherent!

💬 2 🔄 🗑️ ❤️ 12

