

People on the orange site are laughing at this, assuming it's just an ad and that there's nothing to it. Vulnerability researchers I talk to do not think this is a joke. As an erstwhile vuln researcher myself: do not bet against LLMs on this.

[Axios: Anthropic's Claude Opus 4.6 uncovers 500 zero-day flaws in open-source](#)

I think vulnerability research might be THE MOST LLM-amenable software engineering problem. Pattern-driven. Huge corpus of operational public patterns. Closed loops. Forward progress from stimulus/response tooling. Search problems.

Vulnerability research outcomes are in THE MODEL CARDS for frontier labs. Those companies have so much money they're literally distorting the economy. Money buys vuln research outcomes. Why would you think they were faking any of this?

— [Thomas Ptacek](#)

Posted [8th February 2026](#) at 2:25 am

Recent articles

- [How StrongDM's AI team build serious software without even looking at the code](#) - 7th February 2026
- [Running Pydantic's Monty Rust sandboxed Python subset in WebAssembly](#) - 6th February 2026
- [Distributing Go binaries like sqlite-scanner through PyPI using go-to-wheel](#) - 4th February 2026

open-source 285

security 575

thomas-ptacek 16

ai 1824

generative-ai 1616

llms 1581

anthropic 227

claude 226